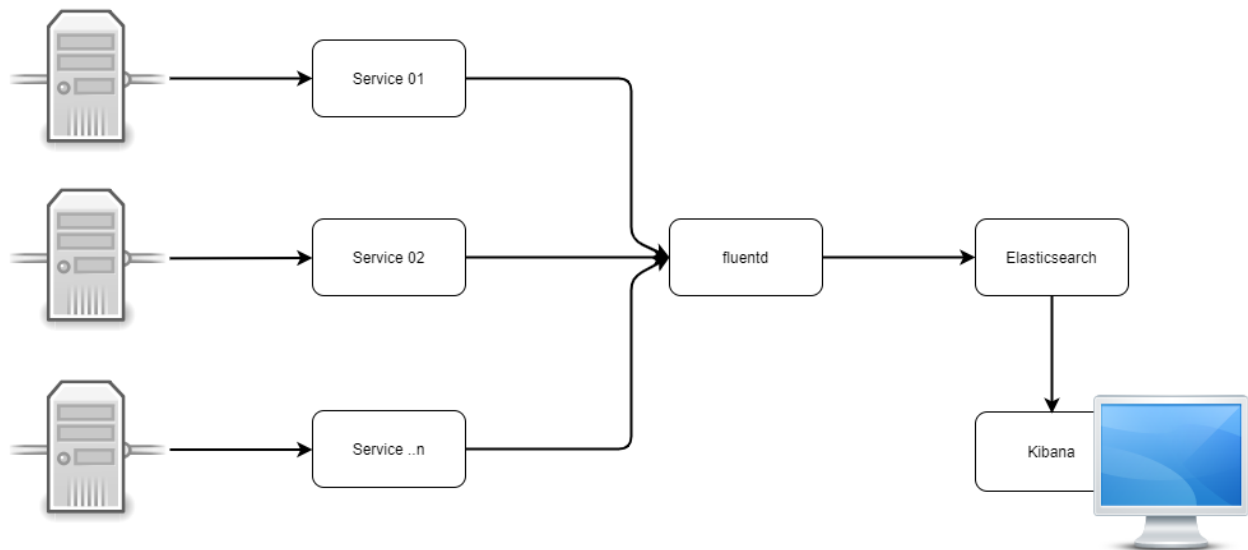


Centralized logging using EFK

Monolithic approach is a default model for creating a software application. A monolithic application is built as a single and indivisible unit which means implementation of logging is similar and only a single log file is enough.

Micro service is an idea to splits application into a set of smaller, interconnected services instead of building a single monolithic application. Every service called container and every container have separate logging. Most difficult part is to handle micro service log without any logging mechanism.

EFK is the solution to centralized logging for micro services and visualize it to users. EFK stand for elasticsearch, fluentd and Kibana.



- **Fluentd** gathers logs from nodes, clean and parse the log data and feeds them to elasticsearch
- **Elasticsearch** is a search and analytics engine for an object store where all logs are stored.
- **Kibana** is a web UI for elasticsearch. lets users visualize data with charts and graphs in Elasticsearch

Getting Started with EFK

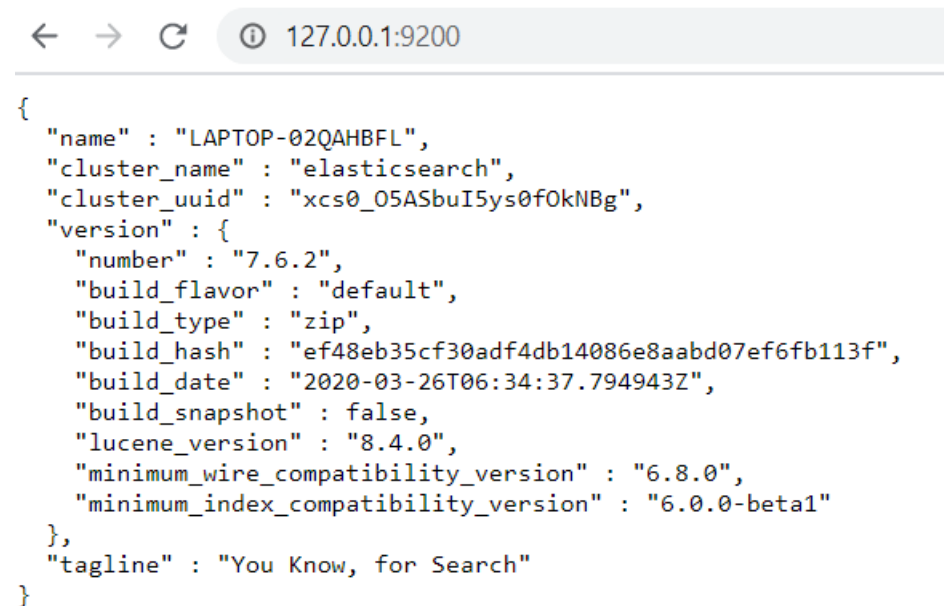
Now have to learn some fundamental concepts of the EFK software that we are going to use, let's get start to deploy on window 10

Step1. Please click the [link](#) to download latest version of elasticsearch for your operating system. unzip folder in a selected drive

Step2. Access bin folder and run elasticsearch.bat in command prompt

```
C:\EFK\elasticsearch-7.6.2\bin>elasticsearch.bat
```

Elasticsearch is running and can be access using url <http://127.0.0.1:9200>



```
{
  "name" : "LAPTOP-02QAHBFL",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "xcs0_05ASbuI5ys0fOkNBg",
  "version" : {
    "number" : "7.6.2",
    "build_flavor" : "default",
    "build_type" : "zip",
    "build_hash" : "ef48eb35cf30adf4db14086e8aabd07ef6fb113f",
    "build_date" : "2020-03-26T06:34:37.794943Z",
    "build_snapshot" : false,
    "lucene_version" : "8.4.0",
    "minimum_wire_compatibility_version" : "6.8.0",
    "minimum_index_compatibility_version" : "6.0.0-beta1"
  },
  "tagline" : "You Know, for Search"
}
```

Step3. Please click [link](#) to download latest version of kibana for your machine. Unzip folder

Step4. Go into kibana-7.6.2\config folder and edit kibana.yml. Uncomment elasticsearch.hosts

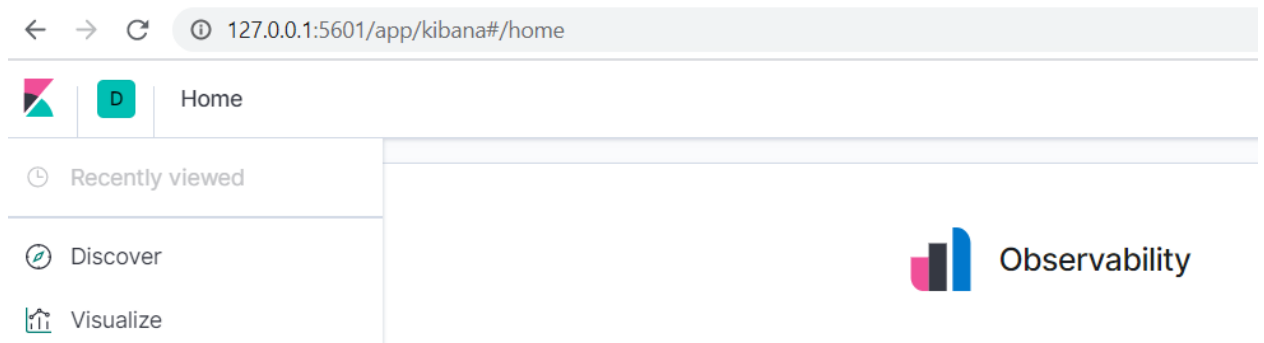
```
# The URLs of the Elasticsearch instances to use for all your queries.
elasticsearch.hosts: ["http://localhost:9200"]
```

Step5. Open command prompt and run kibana.bat

```
C:\EFK\kibana-7.6.2\bin>kibana.bat
```

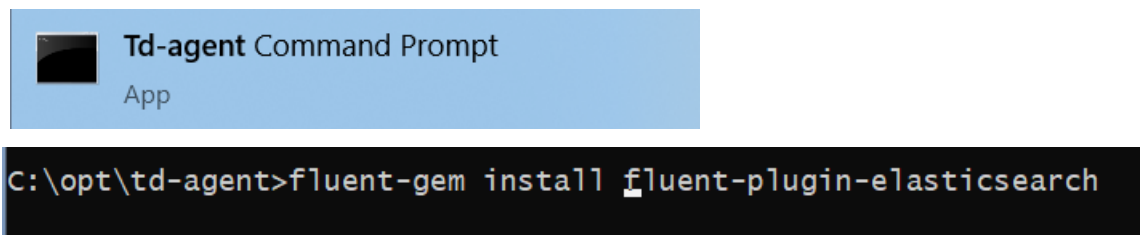
```
log [04:28:07.561] [info][server][Kibana][http] http server running at http://localhost:5601
```

Step6. Kibana running on port 5601 and can be access using url <http://127.0.0.1:5601>



Step7. Please click the [link](#) to download the latest version of fluentd and install it

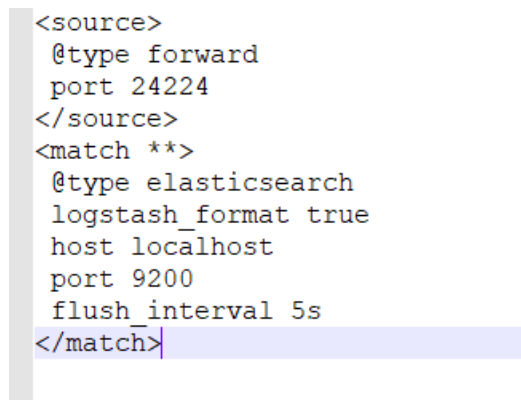
Step8. Open td-agent command prompt and install fluent-plugin-elasticsearch



```
Td-agent Command Prompt
App

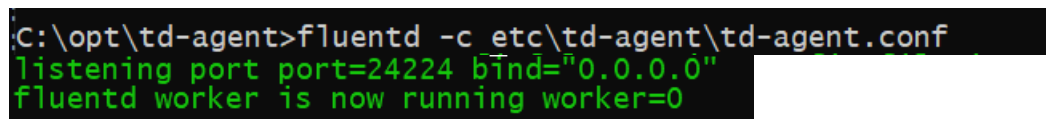
C:\opt\td-agent>fluent-gem install fluent-plugin-elasticsearch
```

Step9. Access folder C:\opt\td-agent\etc\td-agent and open td-agent.conf into editor and modify td-agent.conf file as per below code



```
<source>
  @type forward
  port 24224
</source>
<match **>
  @type elasticsearch
  logstash_format true
  host localhost
  port 9200
  flush_interval 5s
</match>
```

Step10. In the prompt, please execute the command below to launch td-agent process



```
C:\opt\td-agent>fluentd -c etc\td-agent\td-agent.conf
listening port port=24224 bind="0.0.0.0"
fluentd worker is now running worker=0
```

Step11. Open kibana using url <http://localhost:5601> and create new index pattern

← → ↺ localhost:5601/app/kibana#/management/kibana/index_pattern?_g=()

Management / Index patterns / Create index pattern

Elasticsearch

- Index Management
- Index Lifecycle Policies
- Rollup Jobs
- Transforms
- Remote Clusters
- Snapshot and Restore
- License Management
- 8.0 Upgrade Assistant

Kibana

- [Index Patterns](#)

Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for 1

Step 1 of 2: Define index pattern

Index pattern

You can use a * as a wildcard in your index pattern.
You can't use spaces or the characters \, /, ?, ", <, >, |.

Your index pattern can match any of your **4 indices**, below.

Step12. Please clone project folder centralized-logging-using-efk and run springboot application

Step13. Call api inside postman to generate log

GET http://localhost:8086/api/efk/greeting Send

Params Authorization ● Headers (10) Body Pre-request Script Tests Settings

Query Params

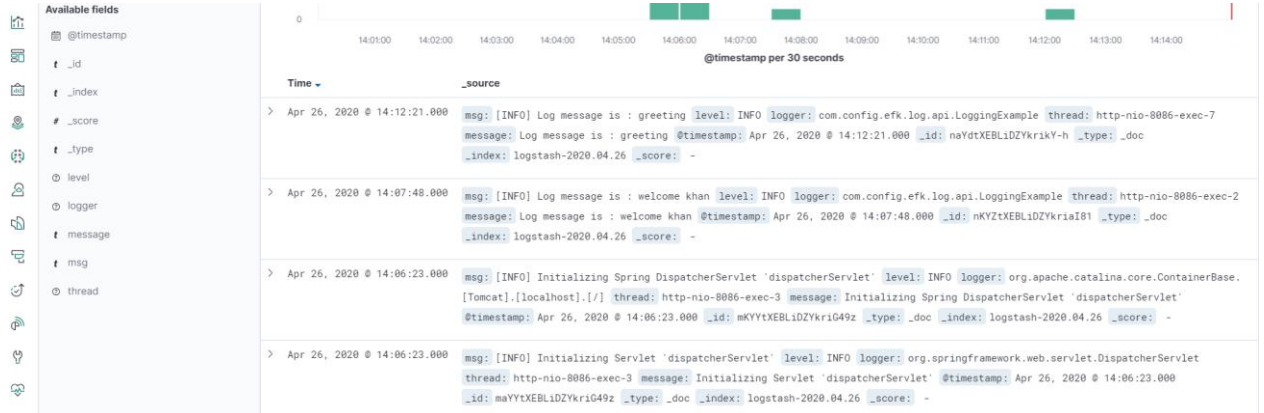
KEY	VALUE	DESCRIPTION
Key	Value	Description

Body Cookies Headers (5) Test Results Status: 200 OK Time: 15 ms Size: 163 B Save

Pretty Raw Preview Visualize JSON

```
1 greeting
```

Step14. Please access kibana dashboard and verify log



Conclusion

This article explains a basic implementation of EFK using single node on local machine. EFK know as platform services. It's better to install EFK separate from application. My organization is using rancher configured with fluentd to send kubernetes logs on elasticsearch